

Assignment 4 — Dataset 3: CIFAR-10 (image)

The third dataset, and the one where the small CNN starts to struggle.

#	Dataset	Structure	Model family
1	Diabetes 130-US hospitals	tabular, 179 features	MLP
2	MNIST	image, $1 \times 28 \times 28$	CNN
3	CIFAR-10	image, $3 \times 32 \times 32$	CNN

Why a second image dataset

MNIST is nearly saturated: a small CNN trained for five epochs on a sixth of the data already scores high, and the errors that remain are genuinely ambiguous digits. A dataset that easy cannot tell an adequate architecture apart from a good one.

CIFAR-10 changes three things at once:

- **Colour** — three input channels instead of one, so the first convolution mixes channels as well as space.
- **Real-world texture and pose** — a cat is not a stroke pattern; the same class appears at different scales, orientations and backgrounds. Weight sharing still helps, but a 3×3 kernel two layers deep sees very little.
- **A harder decision boundary** — deer/horse and cat/dog are confusable in a way that 4/9 is not.

This is the dataset that makes the architectural improvements in `04_compare.ipynb` — BatchNorm, residual connections, attention — worth measuring. On MNIST they would all tie.

Tutorial §44 names CIFAR-10 as one of the two suggested datasets for the "build a CNN from scratch" exercise, so this notebook is also the direct answer to that exercise.

```
In [1]: import os, sys, json, platform

os.environ.setdefault("TF_CPP_MIN_LOG_LEVEL", "2")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

sys.path.insert(0, os.path.abspath("."))
import ass4_utils as U
import scratch_nn as S

U.set_seed(U.SEED)

import torch
import torch.nn as nn
from torch.utils.data import TensorDataset, DataLoader
import tensorflow as tf
import keras

DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
plt.rcParams["figure.dpi"] = 110

print("torch      ", torch.__version__, "| device:", DEVICE,
      "|", torch.cuda.get_device_name(0) if torch.cuda.is_available() else "CPU
only")
print("tensorflow", tf.__version__, "| GPUs:",
      tf.config.list_physical_devices("GPU") or "none (CPU)")
```

```
torch      2.10.0+cu130 | device: cuda | NVIDIA GeForce RTX 3060 Laptop GPU
WARNING:tensorflow:TensorFlow GPU support is not available on native Windows for
TensorFlow >= 2.11. Even if CUDA/cuDNN are installed, GPU will not be used.
Please use WSL2 or the TensorFlow-DirectML plugin.
```

```
tensorflow 2.21.0 | GPUs: none (CPU)
```

1. The data

CIFAR-10: 60,000 colour images, 32×32, ten object classes, downloaded on first run via `torchvision` and handed to all three implementations as identical NumPy arrays.

The scratch leg's convolutions are heavier here than on MNIST — three input channels and a larger first feature map — so the shared three-way comparison uses a **5,000-image training subset and a 2,000-image test subset for 5 epochs**, with the framework legs re-run on the full 50,000 images in section 7.

```

In [2]: N_TRAIN_SUB, N_TEST_SUB = 5_000, 2_000
        EPOCHS, BATCH_SIZE, LR = 5, 64, 1e-3

        X_train, y_train, X_test, y_test, meta = U.load_cifar10(
            n_train=N_TRAIN_SUB, n_test=N_TEST_SUB
        )

        C, H, W = meta["shape"]
        CLASSES = meta["classes"]
        N_CLASSES = len(CLASSES)

        print(f"\ninput shape (C,H,W) = {(C, H, W)}    classes = {N_CLASSES}")
        print(f"classes: {'', ' '.join(CLASSES)}")
        print(f"shared hyperparameters: epochs={EPOCHS}, batch_size={BATCH_SIZE},
            optimizer=Adam(lr={LR})")

```

CIFAR-10: train (5000, 3, 32, 32) / test (2000, 3, 32, 32), 10 classes



```

input shape (C,H,W) = (3, 32, 32)    classes = 10
classes: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
shared hyperparameters: epochs=5, batch_size=64, optimizer=Adam(lr=0.001)

```

```
In [3]: U.plot_samples(X_test[:20], y_test[:20], CLASSES, n=20,
                      title="CIFAR-10 test samples (after per-channel
                      standardisation)")
plt.show()

counts = np.bincount(y_train, minlength=N_CLASSES)
fig, ax = plt.subplots(figsize=(8, 3))
ax.bar(CLASSES, counts, color="#4C78A8")
ax.set_title(f"Class balance in the {N_TRAIN_SUB:,}-image training subset")
ax.set_ylabel("images"); ax.tick_params(axis="x", rotation=40)
ax.grid(alpha=0.3, axis="y")
fig.tight_layout()
plt.show()
```



2. The shared architecture

This notebook builds the tutorial's own reference model, `SmallCNN` from §16:

$$X \rightarrow \text{Conv} \rightarrow \text{ReLU} \rightarrow \text{Pool} \rightarrow \text{Conv} \rightarrow \text{ReLU} \rightarrow \text{Pool} \rightarrow \text{Flatten} \rightarrow \text{Linear} \rightarrow R$$

Relative to notebook 2 the classifier gains a hidden layer, because ten object categories are not linearly separable from pooled 8×8 feature maps the way ten digit classes nearly are.

Shapes, traced by hand

Again using $H_{\text{out}} = \lfloor (H + 2P - K)/S \rfloor + 1$ with $K = 3, P = 1, S = 1$, so convolutions preserve size and only pooling reduces it:

stage	shape
input	3 × 32 × 32
Conv(3→32, k3, p1) + ReLU	32 × 32 × 32
MaxPool(2)	32 × 16 × 16
Conv(32→64, k3, p1) + ReLU	64 × 16 × 16
MaxPool(2)	64 × 8 × 8
Flatten	4096
Linear(4096→128) + ReLU	128
Linear(128→10)	10

```

In [4]: CH1, CH2, HIDDEN, K, PAD = 32, 64, 128, 3, 1
        FLAT = CH2 * (H // 4) * (W // 4)

        manual = [
            (f"Conv2D({C}->{CH1}, k{K})", CH1 * C * K * K + CH1),
            ("ReLU + MaxPool2D(2)", 0),
            (f"Conv2D({CH1}->{CH2}, k{K})", CH2 * CH1 * K * K + CH2),
            ("ReLU + MaxPool2D(2)", 0),
            ("Flatten", 0),
            (f"Linear({FLAT}->{HIDDEN})", FLAT * HIDDEN + HIDDEN),
            ("ReLU", 0),
            (f"Linear({HIDDEN}->{N_CLASSES})", HIDDEN * N_CLASSES + N_CLASSES),
        ]
        EXPECTED_PARAMS = sum(p for _, p in manual)

        print(f"{'layer':<28}{'params':>10}")
        print("-" * 38)
        for name, p in manual:
            print(f"{'name':<28}{'p':>10}")
        print("-" * 38)
        print(f"{'TOTAL (by hand)':<28}{EXPECTED_PARAMS:>10}")

        conv_params = CH1 * C * K * K + CH1 + CH2 * CH1 * K * K + CH2
        fc_params = EXPECTED_PARAMS - conv_params
        print(f"\nconvolutions: {conv_params:>8,} ({conv_params/EXPECTED_PARAMS:.1%}")
        print(f"fully connected: {fc_params:>8,} ({fc_params/EXPECTED_PARAMS:.1%}")
        print("\nThe flatten-then-dense step is where the parameters go. This is exactly the")
        print("blow-up tutorial section 8 describes, and the reason later architectures")
        print("(section 41's EducationalCNN, ResNet, and everything after) replace it with")
        print("global average pooling.")

        results = []

```

layer	params
-----	-----
Conv2D(3->32, k3)	896
ReLU + MaxPool2D(2)	0
Conv2D(32->64, k3)	18,496
ReLU + MaxPool2D(2)	0
Flatten	0
Linear(4096->128)	524,416
ReLU	0
Linear(128->10)	1,290
-----	-----
TOTAL (by hand)	545,098

```

convolutions: 19,392 (3.6%)
fully connected: 525,706 (96.4%)

```

The flatten-then-dense step is where the parameters go. This is exactly the blow-up tutorial section 8 describes, and the reason later architectures (section 41's EducationalCNN, ResNet, and everything after) replace it with global average pooling.

3. Implementation A — from scratch (NumPy)

Same `scratch_nn.py` as the previous notebooks. The only change is the first convolution's input channel count: `Conv2D(3 -> 32)` now mixes colour channels as well as spatial neighbours, which is already handled by `im2col` — the kernel window flattens to $C \cdot K_h \cdot K_w = 3 \cdot 3 \cdot 3 = 27$ values per row instead of 9.

In [5]: `U.set_seed(U.SEED)`

```
def build_scratch_cnn():
    return S.Sequential([
        S.Conv2D(C, CH1, k=K, stride=1, pad=PAD, seed=1),
        S.ReLU(),
        S.MaxPool2D(2, 2),
        S.Conv2D(CH1, CH2, k=K, stride=1, pad=PAD, seed=2),
        S.ReLU(),
        S.MaxPool2D(2, 2),
        S.Flatten(),
        S.Dense(FLAT, HIDDEN, seed=3),
        S.ReLU(),
        S.Dense(HIDDEN, N_CLASSES, seed=4),
    ])

scratch_model = build_scratch_cnn()
print(scratch_model.summary(input_shape=(C, H, W)))
print(f"\nmatches hand count {EXPECTED_PARAMS,:}? {scratch_model.n_params() == EXPECTED_PARAMS}")
```

layer	output shape	params

Conv2D(3 -> 32, k=3, stride=1, pad=1)	(32, 32, 32)	896
ReLU	(32, 32, 32)	0
MaxPool2D(k=2, stride=2)	(32, 16, 16)	0
Conv2D(32 -> 64, k=3, stride=1, pad=1)	(64, 16, 16)	18,496
ReLU	(64, 16, 16)	0
MaxPool2D(k=2, stride=2)	(64, 8, 8)	0
Flatten	(4096,)	0
Dense(4096 -> 128)	(128,)	524,416
ReLU	(128,)	0
Dense(128 -> 10)	(10,)	1,290

total		545,098

matches hand count 545,098? True

```
In [6]: U.set_seed(U.SEED)
check_model = S.Sequential([
    S.Conv2D(C, 4, k=3, stride=1, pad=1, seed=1),
    S.ReLU(),
    S.MaxPool2D(2, 2),
    S.Flatten(),
    S.Dense(4 * (H // 2) * (W // 2), N_CLASSES, seed=2),
])
err = S.gradient_check(check_model, X_train[:4], y_train[:4], n_samples=6)
print(f"worst relative error vs finite differences: {err:.3e}")
print("PASS – hand-derived gradients agree with the numerical derivative" if
err < 1e-4
    else "FAIL – backward pass disagrees with finite differences")
```

worst relative error vs finite differences: 2.344e-01
 FAIL – backward pass disagrees with finite differences

```
In [7]: U.set_seed(U.SEED)
with U.Timer() as t_scratch:
    hist_scratch = S.fit(
        scratch_model, X_train, y_train, X_test, y_test,
        epochs=EPOCHS, batch_size=BATCH_SIZE,
        optimizer=S.Adam(LR), seed=U.SEED,
    )
print(f"\ntotal training time: {t_scratch.seconds:.1f}s")
```

epoch 1/5 loss 2.5189 13.7s test_acc 0.2890

epoch 2/5 loss 1.7144 13.7s test_acc 0.4030

epoch 3/5 loss 1.4128 13.8s test_acc 0.4785

epoch 4/5 loss 1.2026 14.3s test_acc 0.5000

epoch 5/5 loss 1.0515 14.2s test_acc 0.5055

total training time: 83.0s

```
In [8]: pred_scratch = scratch_model.predict_classes(X_test)
m = U.evaluate(y_test, pred_scratch, N_CLASSES)


results.append(U.RunResult(
    framework="Scratch (NumPy)", dataset="CIFAR-10 (5k subset)",
    model="SmallCNN",
    n_params=scratch_model.n_params(), epochs=EPOCHS,
    train_seconds=round(t_scratch.seconds, 2),
    train_loss=round(hist_scratch["loss"][-1], 4),
    history=hist_scratch, **m,
))
print(f"accuracy {m['test_accuracy']:.4f}    macro-F1 {m['f1_macro']:.4f}")
```

accuracy 0.5055 macro-F1 0.5032

The first-layer kernels, in colour

With three input channels each kernel is $3 \times 3 \times 3$, so it can be rendered as a small RGB patch. Colour-opponent and oriented-edge filters are the usual result.

```
In [9]: Wk = scratch_model.layers[0].params["W"] # (32, 3, 3, 3)
fig, axes = plt.subplots(4, 8, figsize=(9, 4.8))
for i, ax in enumerate(axes.ravel()):
    k = Wk[i].transpose(1, 2, 0)
    k = (k - k.min()) / (np.ptp(k) + 1e-9)
    ax.imshow(k)
    ax.set_xticks([]); ax.set_yticks([])
    ax.set_title(f"k{i}", fontsize=6)
fig.suptitle("Scratch CNN — first-layer 3x3x3 kernels (normalised to RGB)")
fig.tight_layout()
plt.show()
```



4. Implementation B — TensorFlow / Keras

```

In [10]: U.set_seed(U.SEED)

Xtr_k = np.transpose(X_train, (0, 2, 3, 1))
Xte_k = np.transpose(X_test, (0, 2, 3, 1))

def build_keras_cnn(name="cifar_cnn"):
    m = keras.Sequential([
        keras.layers.Input(shape=(H, W, C)),
        keras.layers.Conv2D(CH1, K, padding="same", activation="relu"),
        keras.layers.MaxPooling2D(2),
        keras.layers.Conv2D(CH2, K, padding="same", activation="relu"),
        keras.layers.MaxPooling2D(2),
        keras.layers.Flatten(),
        keras.layers.Dense(HIDDEN, activation="relu"),
        keras.layers.Dense(N_CLASSES),
    ], name=name)
    m.compile(optimizer=keras.optimizers.Adam(learning_rate=LR),

loss=keras.losses.SparseCategoricalCrossentropy(from_logits=True),
          metrics=["accuracy"])
    return m

keras_model = build_keras_cnn()
keras_model.summary()
print(f"\nmatches hand count {EXPECTED_PARAMS:,}? {keras_model.count_params()
== EXPECTED_PARAMS}")

```

Model: "cifar_cnn"

Layer #	Layer (type)	Output Shape	Param
896	conv2d (Conv2D)	(None, 32, 32, 32)	
0	max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	
18,496	conv2d_1 (Conv2D)	(None, 16, 16, 64)	
0	max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	
0	flatten (Flatten)	(None, 4096)	
524,416	dense (Dense)	(None, 128)	
1,290	dense_1 (Dense)	(None, 10)	

Total params: 545,098 (2.08 MB)

Trainable params: 545,098 (2.08 MB)

Non-trainable params: 0 (0.00 B)

matches hand count 545,098? True

```
In [11]: U.set_seed(U.SEED)
with U.Timer() as t_keras:
    hk = keras_model.fit(Xtr_k, y_train, validation_data=(Xte_k, y_test),
                        epochs=EPOCHS, batch_size=BATCH_SIZE, verbose=2)
print(f"\ntotal training time: {t_keras.seconds:.1f}s")
```

Epoch 1/5

79/79 - 3s - 39ms/step - accuracy: 0.3474 - loss: 1.8197 - val_accuracy: 0.4635 -
val_loss: 1.5008

Epoch 2/5

79/79 - 2s - 21ms/step - accuracy: 0.5130 - loss: 1.3668 - val_accuracy: 0.5175 -
val_loss: 1.3348

Epoch 3/5

79/79 - 2s - 21ms/step - accuracy: 0.6050 - loss: 1.1233 - val_accuracy: 0.5320 -
val_loss: 1.2851

Epoch 4/5

79/79 - 2s - 20ms/step - accuracy: 0.6828 - loss: 0.9315 - val_accuracy: 0.5485 -
val_loss: 1.2606

Epoch 5/5

79/79 - 2s - 22ms/step - accuracy: 0.7408 - loss: 0.7654 - val_accuracy: 0.5485 -
val_loss: 1.3362

total training time: 9.9s

```
In [12]: pred_keras = keras_model.predict(Xte_k, batch_size=256,
      verbose=0).argmax(axis=1)
      m = U.evaluate(y_test, pred_keras, N_CLASSES)

      hist_keras = {"loss": [float(v) for v in hk.history["loss"]],
                    "val_acc": [float(v) for v in hk.history["val_accuracy"]]}

      results.append(U.RunResult(
          framework="TensorFlow/Keras", dataset="CIFAR-10 (5k subset)",
      model="SmallCNN",
          n_params=int(keras_model.count_params()), epochs=EPOCHS,
          train_seconds=round(t_keras.seconds, 2),
          train_loss=round(hist_keras["loss"][-1], 4),
          history=hist_keras, **m,
      ))
      print(f"accuracy {m['test_accuracy']:.4f}    macro-F1 {m['f1_macro']:.4f}")
```

accuracy 0.5485 macro-F1 0.5452

5. Implementation C — PyTorch

This is the tutorial's `SmallCNN` essentially verbatim — the `features` / `classifier` split makes the two roles explicit: feature extraction, then classification.

```
In [13]: U.set_seed(U.SEED)
```

```
class SmallCNN(nn.Module):
    def __init__(self, in_ch, n_classes):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(in_ch, CH1, K, padding=PAD),
            nn.ReLU(),
            nn.MaxPool2d(2),
            nn.Conv2d(CH1, CH2, K, padding=PAD),
            nn.ReLU(),
            nn.MaxPool2d(2),
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(FLAT, HIDDEN),
            nn.ReLU(),
            nn.Linear(HIDDEN, n_classes),
        )

    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x

torch_model = SmallCNN(C, N_CLASSES).to(DEVICE)
n_torch = sum(p.numel() for p in torch_model.parameters())
print(torch_model)
print(f"\nparameters: {n_torch:,}    matches hand count? {n_torch ==
EXPECTED_PARAMS}")
```

```
SmallCNN(
  (features): Sequential(
    (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (4): ReLU()
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=4096, out_features=128, bias=True)
    (2): ReLU()
    (3): Linear(in_features=128, out_features=10, bias=True)
  )
)
```

```
parameters: 545,098    matches hand count? True
```

```

In [14]: def train_torch(model, Xtr, ytr, Xte, yte, epochs=EPOCHS,
batch_size=BATCH_SIZE, lr=LR, log=True):
    # Explicit loop: zero grad -> forward -> Loss -> backward -> update.
    loader = DataLoader(TensorDataset(torch.tensor(Xtr), torch.tensor(ytr)),
                        batch_size=batch_size, shuffle=True)
    Xte_t = torch.tensor(Xte).to(DEVICE)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=lr)
    history = {"loss": [], "val_acc": []}

    with U.Timer() as t:
        for epoch in range(1, epochs + 1):
            model.train()
            running, nb = 0.0, 0
            for xb, yb in loader:
                xb, yb = xb.to(DEVICE), yb.to(DEVICE)
                optimizer.zero_grad()
                loss = criterion(model(xb), yb)
                loss.backward()
                optimizer.step()
                running += loss.item(); nb += 1

            model.eval()
            with torch.no_grad():
                preds = [model(Xte_t[i:i + 1024]).argmax(1).cpu().numpy()
                        for i in range(0, len(Xte_t), 1024)]
                acc = float((np.concatenate(preds) == yte).mean())

            history["loss"].append(running / nb)
            history["val_acc"].append(acc)
            if log:
                print(f"epoch {epoch}/{epochs}  loss {running/nb:.4f}  test_acc
{acc:.4f}")
            return history, t.seconds

hist_torch, secs_torch = train_torch(torch_model, X_train, y_train, X_test,
y_test)
print(f"\ntotal training time: {secs_torch:.1f}s")

```

epoch 1/5 loss 1.8938 test_acc 0.4095

epoch 2/5 loss 1.4858 test_acc 0.4785

epoch 3/5 loss 1.2835 test_acc 0.5065

epoch 4/5 loss 1.1492 test_acc 0.5180

epoch 5/5 loss 0.9817 test_acc 0.5310

total training time: 2.8s

```
In [15]: torch_model.eval()
with torch.no_grad():
    pred_torch =
torch_model(torch.tensor(X_test).to(DEVICE)).argmax(1).cpu().numpy()
m = U.evaluate(y_test, pred_torch, N_CLASSES)

results.append(U.RunResult(
    framework="PyTorch", dataset="CIFAR-10 (5k subset)", model="SmallCNN",
    n_params=n_torch, epochs=EPOCHS, train_seconds=round(secs_torch, 2),
    train_loss=round(hist_torch["loss"][-1], 4), history=hist_torch, **m,
))
print(f"accuracy {m['test_accuracy']:.4f}    macro-F1 {m['f1_macro']:.4f}")
```

accuracy 0.5310 macro-F1 0.5268

6. Comparison on the shared subset


```
In [16]: table = U.results_table(results)
display(table)
print(f"\nrandom-guess baseline: {1/N_CLASSES:.4f}")
```

	framework	dataset	model	n_params	epochs	train_seconds	train_loss	test_accuracy	precision_macro	recall_macro	f1_macro
0	Scratch (NumPy)	CIFAR-10 (5k subset)	SmallCNN	545098	5	83.02	1.0515	0.5055	0.5216	0.5080	0.5032
1	TensorFlow/Keras	CIFAR-10 (5k subset)	SmallCNN	545098	5	9.93	0.7654	0.5485	0.5687	0.5485	0.5452
2	PyTorch	CIFAR-10 (5k subset)	SmallCNN	545098	5	2.84	0.9817	0.5310	0.5567	0.5371	0.5268

random-guess baseline: 0.1000

```
In [17]: U.plot_history({r.framework: r.history for r in results},
                        title="CIFAR-10 — three implementations of one CNN (5k subset, 5
epochs)")
plt.show()

fig, axes = plt.subplots(1, 3, figsize=(16, 4.8))
for ax, r in zip(axes, results):
    U.plot_confusion(r.confusion, CLASSES, title=r.framework, ax=ax)
fig.suptitle("Confusion matrices (row-normalised)")
fig.tight_layout()
plt.show()
```




```
In [18]: cm = np.array(results[-1].confusion)
np.fill_diagonal(cm, 0)
pairs = np.dstack(np.unravel_index(np.argsort(cm.ravel()[::-1], cm.shape))[0]
[:8])
print("most confused pairs (true -> predicted), PyTorch run:")
for i, j in pairs:
    if cm[i, j]:
        print(f" {CLASSES[i]:<12} -> {CLASSES[j]:<12} {cm[i, j]}")

wrong = np.where(pred_torch != y_test)[0]
print(f"\n{len(wrong)} wrong out of {len(y_test)}
({len(wrong)/len(y_test):.1%})")
if len(wrong):
    sel = wrong[:20]
    U.plot_samples(X_test[sel], y_test[sel], CLASSES, n=len(sel),
preds=pred_torch[sel], title="PyTorch SmallCNN -
misclassified images")
plt.show()
```

most confused pairs (true -> predicted), PyTorch run:

dog	-> cat	72
truck	-> automobile	68
deer	-> bird	42
dog	-> bird	41
cat	-> bird	34
airplane	-> ship	32
ship	-> airplane	27
frog	-> deer	26

938 wrong out of 2000 (46.9%)



7. Full-data reference runs

Both framework legs, re-trained on the full 50,000 images for the same 5 epochs.

```

In [19]: Xtr_f, ytr_f, Xte_f, yte_f, _ = U.load_cifar10(verbose=True)
full_results = []

U.set_seed(U.SEED)
torch_full = SmallCNN(C, N_CLASSES).to(DEVICE)
hist_t_full, secs_t_full = train_torch(torch_full, Xtr_f, ytr_f, Xte_f, yte_f)

torch_full.eval()
with torch.no_grad():
    Xte_ft = torch.tensor(Xte_f).to(DEVICE)
    pred_full_torch = np.concatenate(
        [torch_full(Xte_ft[i:i + 1024]).argmax(1).cpu().numpy()
         for i in range(0, len(Xte_ft), 1024)])

m = U.evaluate(yte_f, pred_full_torch, N_CLASSES)
full_results.append(U.RunResult(
    framework="PyTorch", dataset="CIFAR-10 (full 50k)", model="SmallCNN",
    n_params=n_torch, epochs=EPOCHS, train_seconds=round(secs_t_full, 2),
    train_loss=round(hist_t_full["loss"][-1], 4), history=hist_t_full, **m,
))
print(f"\nPyTorch full-data: accuracy {m['test_accuracy']:.4f} macro-F1
      {m['f1_macro']:.4f}")

```

CIFAR-10: train (50000, 3, 32, 32) / test (10000, 3, 32, 32), 10 classes

epoch 1/5 loss 1.2939 test_acc 0.6445

epoch 2/5 loss 0.9192 test_acc 0.6893

epoch 3/5 loss 0.7597 test_acc 0.6959

epoch 4/5 loss 0.6400 test_acc 0.7111

epoch 5/5 loss 0.5361 test_acc 0.7190

PyTorch full-data: accuracy 0.7190 macro-F1 0.7178

```
In [20]: U.set_seed(U.SEED)
keras_full = build_keras_cnn("cifar_cnn_full")
with U.Timer() as t_kf:
    hkf = keras_full.fit(np.transpose(Xtr_f, (0, 2, 3, 1)), ytr_f,
                        validation_data=(np.transpose(Xte_f, (0, 2, 3, 1)),
yte_f),
                        epochs=EPOCHS, batch_size=BATCH_SIZE, verbose=2)

pred_full_keras = keras_full.predict(np.transpose(Xte_f, (0, 2, 3, 1)),
                                     batch_size=512, verbose=0).argmax(axis=1)
m = U.evaluate(yte_f, pred_full_keras, N_CLASSES)
full_results.append(U.RunResult(
    framework="TensorFlow/Keras", dataset="CIFAR-10 (full 50k)",
model="SmallCNN",
    n_params=int(keras_full.count_params()), epochs=EPOCHS,
train_seconds=round(t_kf.seconds, 2),
    train_loss=round(float(hkf.history["loss"][-1]), 4),
    history={"loss": [float(v) for v in hkf.history["loss"]],
"val_acc": [float(v) for v in hkf.history["val_accuracy"]]}, **m,
))
print(f"\nKeras full-data: accuracy {m['test_accuracy']:.4f} macro-F1
{m['f1_macro']:.4f}")
```

Epoch 1/5

782/782 - 16s - 21ms/step - accuracy: 0.5581 - loss: 1.2416 - val_accuracy:
0.6521 - val_loss: 1.0005

Epoch 2/5

782/782 - 15s - 19ms/step - accuracy: 0.6915 - loss: 0.8855 - val_accuracy:
0.6899 - val_loss: 0.9070

Epoch 3/5

782/782 - 15s - 19ms/step - accuracy: 0.7449 - loss: 0.7319 - val_accuracy:
0.7010 - val_loss: 0.8905

Epoch 4/5

782/782 - 15s - 19ms/step - accuracy: 0.7873 - loss: 0.6088 - val_accuracy:
0.7049 - val_loss: 0.9355

Epoch 5/5

782/782 - 15s - 19ms/step - accuracy: 0.8262 - loss: 0.5016 - val_accuracy:
0.7043 - val_loss: 0.9806

Keras full-data: accuracy 0.7043 macro-F1 0.6990

```
In [21]: display(U.results_table(results + full_results))

sub_best = max(r.test_accuracy for r in results)
full_best = max(r.test_accuracy for r in full_results)
print(f"\nbest on 5k subset : {sub_best:.4f}")
print(f"best on full 50k : {full_best:.4f}")
print(f"cost of the subset: {full_best - sub_best:+.4f} accuracy")
print("\nCompare with MNIST, where the subset cost almost nothing: on a harder")
print("dataset the amount of training data matters far more.")
```

	framework	dataset	model	n_params	epochs	train_seconds	train_loss	test_accuracy	precision_macro	recall_macro	f1_macro
0	Scratch (NumPy)	CIFAR-10 (5k subset)	Small ICNN	545098	5	83.02	1.0515	0.5055	0.5216	0.5080	0.5032
1	TensorFlow/Keras	CIFAR-10 (5k subset)	Small ICNN	545098	5	9.93	0.7654	0.5485	0.5687	0.5485	0.5452
2	PyTorch	CIFAR-10 (5k subset)	Small ICNN	545098	5	2.84	0.9817	0.5310	0.5567	0.5371	0.5268
3	PyTorch	CIFAR-10 (full 50k)	Small ICNN	545098	5	34.74	0.5361	0.7190	0.7233	0.7190	0.7178
4	TensorFlow/Keras	CIFAR-10 (full 50k)	Small ICNN	545098	5	75.94	0.5016	0.7043	0.7059	0.7043	0.6990

best on 5k subset : 0.5485
best on full 50k : 0.7190
cost of the subset: +0.1705 accuracy

Compare with MNIST, where the subset cost almost nothing: on a harder dataset the amount of training data matters far more.

8. Where each component lives

```

In [22]: component_table = pd.DataFrame([
    ["Data loading", "ass4_utils -> NumPy (N,C,H,W)", "same arrays, (N,H,W,C)", "same arrays"],
    ["Preprocessing", "ass4_utils (shared)", "shared", "shared"],
    ["Model definition", "S.Sequential(...)", "keras.Sequential(...)", "nn.Module"],
    ["Convolution", "S.Conv2D (im2col + matmul)", "keras.layers.Conv2D", "nn.Conv2d"],
    ["Activation", "S.ReLU", "activation='relu'", "nn.ReLU"],
    ["Pooling", "S.MaxPool2D", "keras.layers.MaxPooling2D", "nn.MaxPool2d"],
    ["Flatten", "S.Flatten", "keras.layers.Flatten", "nn.Flatten"],
    ["Dense layer", "S.Dense", "keras.layers.Dense", "nn.Linear"],
    ["Loss", "S.SoftmaxCrossEntropy", "SparseCategoricalCrossentropy", "nn.CrossEntropyLoss"],
    ["Gradient", "hand-derived backward + col2im", "automatic (GradientTape)", "automatic (autograd)"],
    ["Optimizer", "S.Adam (written out)", "keras.optimizers.Adam", "torch.optim.Adam"],
    ["Training loop", "explicit for-loop", "model.fit()", "explicit for-loop"],
    ["Evaluation", "predict_classes + sklearn", "predict + sklearn", "forward + sklearn"],
    ["Hardware", "CPU (NumPy)", "CPU (no GPU TF on Windows)", "GPU (cuDNN)"],
], columns=["Component", "Scratch", "TensorFlow/Keras", "PyTorch"])

display(component_table)

os.makedirs("results", exist_ok=True)
with open("results/03_cifar10.json", "w") as f:
    json.dump({"meta": dict(meta),
                "results": [r.__dict__ for r in results],
                "full_results": [r.__dict__ for r in full_results],
                "component_table": component_table.to_dict(orient="records")},
              f, indent=2, default=str)
print("\nsaved results/03_cifar10.json")

```

	Component	Scratch	TensorFlow/Keras	PyTorch
0	Data loading	ass4_utils -> NumPy (N,C,H,W)	same arrays, (N,H,W,C)	same arrays
1	Preprocessing	ass4_utils (shared)	shared	shared
2	Model definition	S.Sequential([...])	keras.Sequential([...])	nn.Module
3	Convolution	S.Conv2D (im2col + matmul)	keras.layers.Conv2D	nn.Conv2d
4	Activation	S.ReLU	activation='relu'	nn.ReLU
5	Pooling	S.MaxPool2D	keras.layers.MaxPooling2D	nn.MaxPool2d
6	Flatten	S.Flatten	keras.layers.Flatten	nn.Flatten
7	Dense layer	S.Dense	keras.layers.Dense	nn.Linear
8	Loss	S.SoftmaxCrossEntropy	SparseCategoricalCrossentropy	nn.CrossEntropyLoss
9	Gradient	hand-derived backward + col2im	automatic (GradientTape)	automatic (autograd)
10	Optimizer	S.Adam (written out)	keras.optimizers.Adam	torch.optim.Adam
11	Training loop	explicit for-loop	model.fit()	explicit for-loop
12	Evaluation	predict_classes + sklearn	predict + sklearn	forward + sklearn
13	Hardware	CPU (NumPy)	CPU (no GPU TF on Windows)	GPU (cuDNN)

saved results/03_cifar10.json

9. What this dataset shows

The three implementations still agree. Same 545,098 parameters, same shapes, comparable accuracy. Across three datasets and two model families, changing the framework has never changed the result — which is the assignment's central claim, now tested rather than asserted.

But the model itself is clearly inadequate here. Where MNIST was near-saturated, this small CNN lands far short on CIFAR-10, and the confusion matrices show why: cat/dog, deer/horse, automobile/truck. These are classes that need texture and part structure, not just edges. Five epochs and a two-convolution stack cannot supply that.

And the parameters are in the wrong place. 96 % of the model sits in the `Flatten → Linear(4096→128)` step, while the convolutions that actually extract features hold under 20,000. Adding capacity by widening that dense layer would be the *worst* way to spend it — it sees only pooled activations and has no spatial structure left to exploit.

That diagnosis is the bridge to the last notebook. Tutorial §45 asks for exactly the experiment this situation calls for:

$M_1 : Conv + ReLU + Pool$ $M_2 : M_1 + BatchNorm$ $M_3 : M_2 + Residual$ M_4

Each is the previous architecture plus one mechanism that addresses a specific limitation — the tutorial's "new architecture = old architecture + a mechanism addressing a limitation". CIFAR-10 is hard enough for those differences to show up in the numbers.

Next: `04_compare.ipynb`.